

NOTES FOR REINFORCEMENT LEARNING

Yao Yao

realyao21002@gmail.com

1 Multi-arm Bandit: only one state

- ϵ -greedy. To find out the best arm by statistical analysis. At each step/round/time, choose the current best arm with probability (w.p.) $(1 - \epsilon)$ (**exploitation**) or randomly choose an arm w.p. ϵ (**exploration**). However, there is always a gap in regret.
- Upper Confidence Bound (UCB) method. To find out the best arm by statistical analysis, too. At each step/round/time, without a probability, choose the largest 'potential' arm that optimize the value function added with a **Chernoff-Hoeffding bound term**. This is to enable the reward gap to approach 0.
- Gradient-bandit method. To decide a distribution for pulling arms. A naive version of policy gradient method.

Elements: K -arm bandit, with mean μ_1, \dots, μ_K , action $A_t \in \mathcal{A} = \{1, \dots, K\}$ and R_t . At each time step compute estimation of mean - using rewards accumulated **before time T**

$$Q_T(a) = \frac{\sum_{t=1}^{T-1} R_t \mathbf{1}_{A_t=a}}{N_T(a)}, \text{ where } N_T(a) = \sum_{t=1}^{T-1} \mathbf{1}_{A_t=a}.$$

1.1 ϵ -greedy

Randomly pick $A_t \sim U\{1, \dots, K\}$ w.p. $\epsilon \in (0, 1)$, or pick $A_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_T(a)$ w.p. $1 - \epsilon$.

In application - generate an R.V. $\sim U(0, 1)$ to ensure the exploration rate is $P(R.V. \leq \epsilon) = \epsilon$.

: if $R.V. \leq \epsilon$, randomly choose an arm.

There is an iteration in Q -

$$Q_{N+1} = \frac{(N-1) \cdot Q_N + R_N}{N} = Q_N + \frac{R_N - Q_N}{N},$$

where $N+1$ is the local time when **an arm** is being estimated, i.e. N is the last time an arm is pulled.

Bernoulli Bandit $r_i \sim \text{Bernoulli}(\mu_i)$

Define an R.V. $Y \sim U(0, 1)$ and

$$X = \begin{cases} 1, & Y < \mu \\ 0, & Y \geq \mu \end{cases}$$

to ensure the reward of a Bernoulli arm is $r_i = 1$ w.p. μ_i

$$: P(X = 1) = P(Y < \mu) = \mu, P(X = 0) = P(Y \geq \mu) = 1 - \mu.$$

Hyper Param Analysis

The larger ϵ , the more exploration, the faster it converges, but not necessarily converge to good solutions. The more exploitation, the more cumulative rewards, but the slower it converges.

Empirical good choice $\epsilon = 0.1$ or $\epsilon = 0.05$.




Figure 1: ϵ -greedy

1.2 Regret analysis: ϵ -greedy results in a gap

A measure of 'expected loss' compared to always pulling the optimal arm.

$$\text{Regret}(T) = T\mu^* - E[\sum_{k=1}^T R_{A_k}] = T\mu^* - E[\sum_{k=1}^T \mu_{A_k}] = T\mu^* - E[\sum_{a \in \mathcal{A}} N(a)\mu_a],$$

where $\mu^* = \max_i \{\mu_i\}$, μ_{A_k} = estimated mean for arm A_k = some a , $N(a)$ = #times arm a is pulled.

If regret $\rightarrow 0$, then

$$\begin{aligned} \text{Regret} &= T\mu^* - E[\sum_{k=1}^T \mu_{A_k}] \rightarrow 0, \\ \text{AvgReward}(T) &= \frac{E[\sum_{k=1}^T \mu_{A_k}]}{T} \rightarrow \mu^*, \\ \frac{\text{Regret}(T)}{T} &= \mu^* - \text{AvgReward}(T) \rightarrow 0. \end{aligned}$$

In ϵ -greedy, assuming arm 10 as the true best

$$\begin{aligned} \text{Regret} &= T\mu_{10} - \frac{9\epsilon}{10}T\bar{\mu}_{\text{other}} - \frac{10-9\epsilon}{10}T\mu_{10} = \frac{9\epsilon T}{10}(\mu_{10} - \bar{\mu}_{\text{other}}), \\ \text{AvgRegret}(T) &= \frac{\text{Regret}(T)}{T} = \frac{9\epsilon}{10}(\mu_{10} - \bar{\mu}_{\text{other}}) > 0. \end{aligned}$$

There is a non-diminishing gap.

A Naive Approach

Introducing **Chernoff-Hoeffding bound**:

$$P(|Q_n(a) - \mu_a| \leq \delta) \geq 1 - 2e^{-2n\delta^2},$$

where n is the number of times arm a is pulled, μ_a is the true mean reward of arm a .

Then the confidence interval $\mu_a \in [Q_n(a) - \delta, Q_n(a) + \delta]$ is what we focus on. Suppose the gap between the true best and the true second best is known as Δ . Let $1 - 2e^{-2n(\Delta/3)^2} \geq 1 - p$. then $n \geq \frac{9\ln(2/p)}{2\Delta^2}$ would satisfies

$$\forall a \neq 10, Q_n(a) \leq \mu_a + \Delta/3 < \mu_{10} - \Delta/3 \leq Q_n(10) \text{ w.p. at least } 1 - p,$$

to ensure the true best arm 10 has the highest $Q_n(10)$ w.p. $1 - 10p$. There are 10 inequalities, where any of them does not hold w.p.

$$P(\text{any does not hold}) = P(\{\text{1st does not hold}\} \cup \dots \cup \{\text{10th does not hold}\}) \leq 10p.$$

Then the approach is **exploration-then-exploitation** - pull each time for n times, and then always choose the arm with best estimated reward for $T - nk$ times. Note that

$$\text{let } p = 1/T, \text{ then } n = \lceil \frac{9 \ln(2T)}{2\Delta^2} \rceil.$$

Such approach has drawbacks that the user needs to know T as the number of trials, which also makes the exploration and exploitation completely separated and lack flexibility, and Δ as the gap.

1.3 The UCB Approach

This approach helps the Chernoff-Hoeffding bound inequality holds, to make the

$$\frac{\text{Regret}(T)}{T} = \mu^* - \text{AvgReward}(T) = O(1/\sqrt{T}) \rightarrow 0.$$

At each time step, choose an arm to maximize the **Upper Confidence Bound** of the previous confidence interval. Here the selected $\delta_t(a)$ varies at each time across arms, and is $\sqrt{\frac{2 \ln t}{N_t(a)}}$ such that

$$1 - 2e^{-2n\delta_t^2} = 1 - 1/t^4.$$

The larger the time t grows, the higher the probability that $Q_t(a) + \delta_t(a)$ provides a true upper

The **UCB** algorithm:

The UCB algorithm

Initialize $t = 1$, for $a = 1$ to K :

$R \leftarrow \text{bandit}(a)$

$Q(a) \leftarrow R, \quad N(a) \leftarrow 1.$

Repeat forever:

$t \leftarrow t + 1$ maximize UCB

$A \leftarrow \underset{a}{\text{argmax}} \left\{ Q(a) + c \sqrt{\frac{\ln t}{N(a)}} \right\}$

$R \leftarrow \text{bandit}(A), \quad N(A) \leftarrow N(A) + 1$

$Q(A) \leftarrow Q(A) + \frac{1}{N(A)}(R - Q(A))$

c is to be chosen

$c = \sqrt{2}$

In each iteration, instead of picking the arm with largest sample average reward $Q_t(a)$, UCB algorithm **optimistically** picks the one with largest "potential"

$$\mu_a \leq Q_t(a) + \sqrt{\frac{2 \ln t}{N_t(a)}}$$

Figure 2: The UCB approach

bound for the true mean μ_a , although δ_t is changing through time.

The larger the pull number $N_t(a)$ grows, the tighter the upper bound $Q_t(a) + \delta_t(a)$ is compared with the true mean μ_a . Because $\delta_t(a)$ would decrease, letting the upper bound be closer to μ_a .

In each iteration, instead of picking the arm with largest $Q_t(a)$, UCB alg optimistically picks the one with **the largest 'potential'**.

- Every arm will be pulled infinitely. $\delta_t(a)$ gets larger even if arm a is never pulled, then arm a will be pulled until its ucb becomes the largest.

- Pulled suboptimal arm will be played less frequently, because $N_t(a)$ gives a penalty in $\delta_t(a)$.
- Overestimate rarely pulled arms. In particular, if $N_t(a) = 0$, then its $ucb = \infty$, making the unpulled arm pulled first. (However, in practice, we always initialize the first pulling of each arm.)

1.4 Gradient-Bandit

Direct optimization to optimize the distribution for pulling each arm. A naive version of the policy gradient method.

$$\begin{aligned} \min_{\pi} & -\sum_{a=1}^K \mu_a \pi_a \\ \text{s.t. } & \pi \geq 0 \\ & \sum_a \pi_a = 1 \end{aligned}$$

Where, the objective is the negative expected reward when pulling arms according to the distribution π . To ensure $\pi_a \in [0, 1]$, introduce a parameter θ_a for each arm and let

$$\pi_a(\theta) = \frac{\exp(\theta_a)}{\sum_{b=1}^K \exp(\theta_b)}.$$

The objective becomes $\min_{\theta} \{f(\theta) : -\sum_{a=1}^K \mu_a \pi_a(\theta)\}$.

In such method, we update the distribution using gradient descent, which utilize the unknown μ_a . Although this is unknown, we can still have an unbiased estimator of the gradient without knowing μ_a . Fortunately,

$$\frac{\partial f(\theta)}{\partial \theta_a} = \sum_{b \neq a}^K \mu_b \pi_b(\theta) (\pi_b(\theta) - 0) + \mu_a \pi_a(\theta) (\pi_a(\theta) - 1) = E_{b \sim \pi(\theta), R = \text{bandit}(b)} [R \cdot (\pi_a(\theta) - \mathbf{1}_{b=a})].$$

The estimator means that, when you pull an arm b , calculate the partial derivative using the expectation for each arm a .

Then, conduct gradient descent step for each $\theta_a^{t+1} = \theta_a^t - \eta_t R_b (\pi_a^t - \mathbf{1}_{a=b})$, where π_a^t is computed at the beginning of each iteration.

Gradient-bandit algorithm

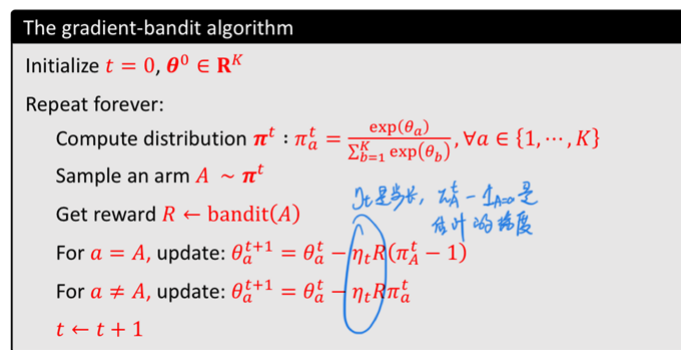


Figure 3: Gradient Bandit

- Taking smaller steps will converge to higher final average reward, but with slower convergence speed.
- Better than ϵ -greedy.

Add base-line to improve GB

The base-line is

$$\bar{R}_T = \frac{\sum_{t=1}^T R_t}{T}, \text{ to update baseline, } \bar{R}_{T+1} = \frac{\bar{R}_T \cdot T + R_{T+1}}{T+1} = \bar{R}_T + \frac{R_{T+1} - \bar{R}_T}{T+1}.$$

When conducting gradient descent,

$$\theta_a^{t+1} = \theta_a^t - \eta_t (R_b - \bar{R}_{t+1}) (\pi_a^t - \mathbf{1}_{a=b}), \text{ when choosing arm } b.$$

Adding a baseline **does not change the unbiased-ness**.

- Ways to improve gradient-bandit
- We can add a base-line to the gradient formula:




Figure 4: gradient bandit with baseline

$$\nabla f_B(\theta) = \nabla f(\theta), \text{ and } \sum_{a=1}^K \pi_a(\theta) = 1 \Rightarrow \sum_{a=1}^K \frac{\partial \pi_a(\theta)}{\partial \theta_a} = 0.$$

Which means equally shifting the rewards of each arm by a constant B will not change the relative performance of the arms and will not change the gradients.

Adding an appropriate base-line will often improve the (overall) performance.

2 Markov Decision Process

The (finite) Markov Decision Process is written as

$$\mathcal{M}(\mathcal{S}, \mathcal{A}, r, \gamma, P).$$

2.1 MDP Elements

- State space \mathcal{S} s.t. $s \in \mathcal{S}$ is a set.
- Action space \mathcal{A} s.t. $a \in \mathcal{A}$. To distinguish between the action choice a and the real action \hat{a} .
- Reward function $r(\cdot, \cdot)$. Results from $s \in \mathcal{S}$ and $a \in \mathcal{A}$, $E[R] = r(s, a)$ or simplified $R = r(s, a)$.
- Transition probability P . A probability $P(s'|s, a)$, so $s' \sim P(\cdot|s, a)$. For each action a , we can also write $P_a(s, s')$, so $s' \sim P_a(s, \cdot)$.

- Discount factor $\gamma \in (0, 1)$. We generally **don't expect a terminating state**, then a discount on future reward should be conducted. We would like to maximize the **expected discounted cumulative reward**

$$E[\sum_{t=0}^{\infty} \gamma^t R_t] \text{ with } E[R_t | s_t, a_t] = r(s_t, a_t).$$

In most cases, γ is close to 1, e.g., $\gamma = 0.95$.

- Policy $\pi : \mathcal{S} \rightarrow \Delta_{\mathcal{A}}$ is a mapping from the state space to a **distribution over the action space**. Under a policy π , the probability of an action a given a state s gives $a \sim \pi(\cdot | s)$. Given an **initial state distribution** ξ , we would like to find the optimal policy by maximizing

$$V^{\pi}(\xi) := E[\sum_{t=0}^{\infty} \gamma^t \cdot r(s_t, a_t) | s_0 \sim \xi, a_t \sim \pi(\cdot | s_t), s_{t+1} \sim P(\cdot | s_t, a_t)].$$

2.2 Estimating Cumulative Reward: Monte Carlo

Approach 1: generate random horizon, without computing γ

At each epoch, generate a random horizon H s.t. $H = X - 1$ with $X \sim \text{Geom}(1 - \gamma)$, then $P(H \geq k) = \gamma^k$. After that, sample a finite sequence $\{s_t, a_t, R_t\}_{t=0}^H$ and compute $\tilde{V}_{epoch} = \sum_{t=0}^H R_t$. After epochs, take the average $\tilde{V} = \frac{1}{h} \sum_{epoch=1}^h \tilde{V}_{epoch}$. The estimator \tilde{V} is unbiased,

$$E[\tilde{V}] = E[\sum_{t=0}^{\infty} \mathbf{1}_{t \leq H} R_t] = \sum_{t=0}^{\infty} P(H \geq t) E[R_t] = \sum_{t=0}^{\infty} \gamma^t E[R_t] = V^{\pi}(\xi).$$

Expected trajectory length is $E[H] + 1 = \frac{1}{1-\gamma}$.

Alternative estimator: fixed horizon

Take a fixed horizon $H = \lceil \text{const} \cdot \frac{\ln(1/\epsilon)}{1-\gamma} \rceil$. Then sample a sequence $\{s_t, a_t, R_t\}_{t=0}^H$ and compute $\tilde{V} = \sum_{t=0}^H \gamma^t R_t$.

This estimator is biased, but the bias is very small:

$$|E[\tilde{V}] - V^{\pi}(\xi)| \leq P(\epsilon).$$

Complimentary knowledge on horizon H

There are two types of horizon, relatively for finite and infinite MDP:

$$H = \begin{cases} \text{terminating state,} & \text{finite MDP} \\ \tau, & \text{infinite MDP} \end{cases}.$$

Fixed horizon may lead to an unexpected terminating state.

3 Value Functions

3.1 State Value Function

Definition

As the cumulative reward function is

$$V^{\pi}(\xi) = E_{s_0 \sim \xi, a_t \sim \pi(\cdot | s_t), s_{t+1} \sim P(\cdot | s_t, a_t)} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)],$$

when the initial state s_0 is known and instead denoted by s , we can compute the state value function from such initial s .

$$V^\pi(s) = E_{a_t \sim \pi(\cdot|s_t), s_{t+1} \sim P(\cdot|s_t, a_t)} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s].$$

Note that

$$V^\pi(\xi) = E_{s_0 \sim \xi} [V^\pi(s_0)].$$

Bellman Equation

$$\begin{aligned} V^\pi(s) &= \sum_{a \in A} \pi(a|s) \sum_{s' \in S} P(s'|s, a) [r(s, a) + \gamma V^\pi(s')] \\ &= \sum_{a \in A} \pi(a|s) [r(s, a) + \sum_{s' \in S} P(s'|s, a) \gamma V^\pi(s')], \forall s. \end{aligned}$$

In expectation form,

$$V^\pi(s) = E_{a \sim \pi(\cdot|s), s' \sim P(\cdot|s, a)} [r(s, a) + \gamma V^\pi(s')], \forall s.$$

$V^\pi(s)$ in Terms of $Q^\pi(s, a)$

$$V^\pi(s) = \sum_{a \in A} \pi(a|s) Q^\pi(s, a) = E_{a \sim \pi(\cdot|s)} Q^\pi(s, a).$$

3.2 State Action Value Function (Q Function)

Definition

$$Q^\pi(s, a) := E_{s_{t+1} \sim P(\cdot|s_t, a_t), a_t \sim \pi(\cdot|s_t)} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s, a_0 = a], \forall s, a.$$

Bellman Equation

$$\begin{aligned} Q^\pi(s, a) &= \sum_{s' \in S} P(s'|s, a) \sum_{a' \in A} \pi(a'|s') [r(s, a) + \gamma Q^\pi(s', a')] \\ &= r(s, a) + \sum_{s' \in S} \sum_{a' \in A} P(s'|s, a) \pi(a'|s') \gamma Q^\pi(s', a'). \end{aligned}$$

In expectation form,

$$Q^\pi(s, a) = r(s, a) + \gamma E_{s' \sim P(\cdot|s, a), a' \sim \pi(\cdot|s')} [Q^\pi(s', a')].$$

$Q^\pi(s, a)$ in Terms of $V^\pi(s)$

$$\begin{aligned} Q^\pi(s, a) &= \sum_{s' \in S} P(s'|s, a) [r(s, a) + \gamma V^\pi(s')] \\ &= r(s, a) + \gamma E_{s' \sim P(\cdot|s, a)} [V^\pi(s')]. \end{aligned}$$

3.3 Optimal Value Function

To obtain the optimal policy π^* , maximize the cumulative reward function

$$\max_{\pi} V^\pi(\xi) := E_{\xi, \pi, P} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)].$$

The obtained π^* will **simultaneously optimize all the value functions and Q functions**.

In such a policy,

$$\pi^*(a|s) = \begin{cases} 1, & a = a^*(s) \\ 0, & a \neq a^*(s) \end{cases}.$$

Where

$$a^* = \operatorname{argmax}_{a \in A} \{Q^*(s, a)\} = \operatorname{argmax}_{a \in A} \{r(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^*(s')\}.$$

Optimal State Value Function

$$V^*(s) = \max_{a \in A} \{r(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V^*(s')\} = \max_{a \in A} \{Q^*(s, a)\}.$$

Optimal State Action Value Function

$$\begin{aligned} Q^*(s, a) &= r(s, a) + \gamma \cdot \max_{a' \in A} \{\sum_{s' \in S} P(s'|s, a) Q^*(s', a')\} \\ &= r(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \cdot \max_{a' \in A} \{Q^*(s', a')\} \end{aligned}$$

3.4 Notes on Solving the Optimal Policy

Note: Solving **Bellman Equations** for $V^*(s)$ or $Q^*(s, a)$ would derive the optimal policy π^* . However, how to solve those Bellman Equations depends on premise assumptions.

- If we assume known transitions and terminating state (i.e. with a horizon), we use dynamic programming, which problem is always called **planning**. This is the focus of section 4.
- If we meet with an infinite-horizon MDP without terminating state, we use value iteration or Q iteration.

4 Finite Horizon MDP: Dynamic Programming

The finite horizon maybe 1) time horizon, or 2) terminating state. (see 2.2)

Note: When setting up, invalid action will be presented by $\text{reward} = -\infty$. In practice, although all the states share the same action set, it is OK for different state to have different action sets.

4.1 Example: Grid World Navigation - Terminating State

- State $s = (i, j)$ as the position.
- Action $a \in \{R, D\}$.
- Transition. The transition $s, a \rightarrow s'$ is deterministic, i.e. $P(s'|s, a) = 1$. Therefore, to simplify, use the transition mapping $s' = \text{trans}(s, a)$. Note that

$$\text{trans}((i, j), R) = (i, \min\{4, j + 1\}),$$

$$\text{trans}((i, j), D) = (\min\{4, i + 1\}, j).$$

Then it can be seen that state $(4, 4)$ is absorbing, which state once hit will not change forever.

- Reward $r(s, a)$ should take the **penalty for invalid state** and take the **terminating state** into consideration.

$$r((i, j), R) = w[\text{trans}((i, j), R)] - \infty \cdot \mathbf{1}_{j=4},$$

$$r((i, j), D) = w[\text{trans}((i, j), D)] - \infty \cdot \mathbf{1}_{i=4}$$

(where w is the reward number on a grid (i, j)). For terminating state,

$$r((4, 4), a) = 0, \forall a \in \{R, D\}.$$

Note: once hit the terminating state, the reward received afterwards would be 0 forever.

Solving Bellman Equations

In such a problem, the probability in the original Bellman Equations would be deterministic as 1 or 0, and we also set $\gamma = 1$ -

$$\begin{aligned} V^*(s) &= \max_{a \in A} \{r(s, a) + \gamma V^*(s')\} \\ &= \max_{a \in A} \{r(s, a) + V^*(\text{trans}(s, a))\}. \end{aligned}$$

In order to conduct **dynamic programming**, start from the terminating state $(4, 4)$ with $V^*(4, 4) = 0$. Then, in the Bellman Equation above, the $r(s, a)$ and the $V^*(s')$ are all known, enabling us to solve the the previous state $V^*(s) = V^*(4, 3)$ and $V^*(s) = V^*(3, 4)$.

Solving those equations yield the optimal policy -

$$\pi^*(a^*|s) = 1, \forall s.$$

Note: for some state, we would yield multiple optimal action a^* , then randomly choose one. In this example, $\pi^*(\cdot|1, 1)$ can be any distribution, for R and D are equally optimal.

4.2 Time Horizon Instead of Terminating State

Horizon of $t = H$ can be view as a special case of MDP with terminating state. **Time can be encoded into the states**, i.e., let $s = (\hat{s}, t) \sim S = \hat{S} \times N$ where $\hat{s} \in S$ and $t \in N$.

In the previous grid world example, state incorporating time is not necessary.

Remind: state incorporating into time information is always by default in many papers, which use π^t to embrace such.

Note: in such cases, all states with $t = H$ are terminating states. Transitions become $P(\cdot|s_t, t, a_t)$. Also, take $\gamma = 1$, then the cumulative reward function will be

$$V^\pi(\xi) = E_{\xi, \pi, P}[\sum_{t=0}^{H-1} r(s_t, a_t)].$$

Maximizing the function above would yield an optimal policy that **depends on time** $a_t \sim \pi(\cdot|s_t, t)$.

Note: whether taking the optimal policy or not,

$$V^\pi(\hat{s}, H) = V^*(\hat{s}, H) = 0, \forall \hat{s} \in \hat{S};$$

$$Q^\pi(\hat{s}, H, a) = Q^*(\hat{s}, H, a) = 0, \forall \hat{s} \in \hat{S}, a \in A.$$

Therefore, the Bellman Equations can be solved **level by level**, i.e., **at each time step**. Knowing

$$V^*(\hat{s}, H-1) = \max_{a \in A} \{r(\hat{s}, a) + \gamma \sum_{\hat{s}' \in \hat{S}} P(\hat{s}'|\hat{s}, H-1, a) V^*(\hat{s}', H)\}.$$

To conduct the **dynamic programming**, also start from the terminating state, i.e., level H - $V^*(\hat{s}, H) = 0, \forall \hat{s} \in \hat{S}$. Then, assuming known reward function $r(\hat{s}, a)$ and transition probability $P(\hat{s}'|\hat{s}, H-1, a)$, as well as the solved $V^*(\hat{s}', H)$, we can solve the previous level $H-1$ states.

4.3 Example: Inventory Management - Time Horizon

- State $s = (n, t)$ as the inventory **at end of** the month t . Assuming no inventory at the beginning of month 1, i.e., $s_0 = (0, 0)$ as the initial state. Therefore, $S = \{(n, t) : n \in \{0, \dots, 4\}, t \in \{1, 2, 3, 4\}\} \cup \{(0, 0)\}$.

- Action $a \in \{0, \dots, 5\}$ as the number of productions **at the beginning of next moth**. To satisfy the **demand and capacity constraints**,

$$a + n \geq d_{t+1}, a + n - d_{t+1} \leq K = 4.$$

- Transition is also **deterministic**. Simply use $s' = \text{trans}(s, a)$ as the transition mapping for $t = 0, 1, 2, 3$ (**not considering transition for $t = 4$**). Let

$$\text{trans}((n, t), a) = (\text{proj}_{[0,4]}(n + a - d_{t+1}), t + 1)$$

- Reward $r((n, t), a)$ in terms of $g = n + a - d_{t+1}$.

$$r((n, t), a) = -0.5 \times \text{Proj}_{[0,4]}(g) - 1 \times a - 3 \times \mathbf{1}_{a>0} - \infty \times \mathbf{1}_{g \notin [0,4]}.$$

to **minimize the total cost**, no revenue is considered, but a penalty for over-capacity inventory is imposed.

- Discount factor $\gamma = 1$ also.
- Time horizon $H = 4$.

4.4 Example: Coin Picking Game - nondeterministic transitions

- State $s = (i, j)$ as the remaining coin from number i to number j . After each step two coins will be picked, therefore the number of remaining coins is even. $S = \{(1, 2), (2, 3), (3, 4), (4, 5), (5, 6), (1, 4), (2, 5), (3, 6), (1, 6)\}$. The terminating state is the null subset O .
- Action $a \in \{L, R\}$.
- Transition is **not deterministic**. Consider from the terminating state:

$$P(O|O, a) = 1, \forall a;$$

$$P(O|(i, i+1), a) = 1, \forall a;$$

$$P((i+2, j)|(i, j), L) = P((i+1, j-1)|(i, j), L) = 0.5,$$

$$P((i, j-2)|(i, j), R) = P((i+1, j-1)|(i, j), R) = 0.5.$$

- Reward is the value of coin picked.

$$r(O, L) = r(O, R) = 0;$$

$$r((i, j), L) = v_i, r((i, j), R) = v_j.$$

The Bellman Equations with $\gamma = 1$

$$V^*(s) = \max_{a \in A} \{r(s, a) + \sum_{s' \in S} P(s'|s, a) V^*(s')\}.$$

To conduct dynamic programming, start from the terminating state -

$$V^*(O) = 0.$$

Then for the previous states, solve all the $V^*(i, i+1)$.

5 Infinite Horizon MDP: Value and Q Iterations

The infinite-horizon MDPs are without any terminating state (**transition dynamics are still known**), and the corresponding approach is value iteration and Q iteration.

Value iteration and Q iteration are iterations using Bellman equations and Bellman operators for

$$V^\pi(\cdot), V^*(\cdot), Q^\pi(\cdot), \text{ and } Q^*(\cdot).$$

Recall that solving V^* and Q^* would derive the optimal policy π^* .

5.1 Fixed-point problem of Bellman operators

Using Bellman operators, we try to solve fixed-point problems including

- Value iteration.

$$V^\pi = T^\pi V^\pi$$

$$V^* = T^* V^*,$$

$$\text{where } T : R^S \rightarrow R^S,$$

and we can write the function for $[T^\pi V^\pi]_s$ and $[T^* V^*]_s$ as

$$[T^\pi V^\pi]_s := \sum_{a \in A} \pi(a|s) [r(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_{s'}^\pi] = V_s^\pi$$

$$[T^* V^*]_s := \max_{a \in A} \{r(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) V_{s'}^*\} = V_s^*.$$

- Q iteration.

$$Q^\pi = T^\pi Q^\pi$$

$$Q^* = T^* Q^*,$$

$$\text{where } T : R^{S \times A} \rightarrow R^{S \times A},$$

and we can write the function for $[T^\pi Q^\pi]_{s,a}$ and $[T^* Q^*]_{s,a}$ as

$$[T^\pi Q^\pi]_{s,a} := r(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \sum_{a' \in A} \pi(a'|s') Q_{s',a'}^\pi = Q_{s,a}^\pi$$

$$[T^* Q^*]_{s,a} := r(s, a) + \gamma \sum_{s' \in S} P(s'|s, a) \max_{a' \in A} \{Q_{s',a'}^*\} = Q_{s,a}^*.$$

5.2 Fixed-point Solution: Contraction Property

Let $0 < \gamma < 1$, the mapping $f(\cdot)$ is a γ -**contraction** under the norm $\|\cdot\|$ if

$$\|f(x) - f(y)\| \leq \gamma \|x - y\|, \forall x, y.$$

Subsequently, the solution to the fixed-point problem $x = f(x)$ can be obtained by the **fixed-point iteration**.

Given the fixed-point problems of value and Q iterations, it can be proved that the Bellman operators are γ -**contraction** under the **infinity** norm, which is defined as $\|x\|_\infty := \max_i |x_i|$. Because they are γ -contraction, the iteration will converge to the true solutions.

$$\|x^t - x^*\| = \|f(x^{t-1}) - f(x^*)\| \leq \gamma \|x^{t-1} - x^*\| \leq \gamma^2 \|x^{t-2} - x^*\| \cdots \leq \gamma^t \|x^0 - x^*\|,$$

$$\Rightarrow \lim_{t \rightarrow \infty} x^t = x^*.$$

Value Iteration

Start from arbitrary an $V^0 \in \mathbf{R}^S$ and set $t = 0$
Based the purpose, set $T = T^\pi$ or $T = T^*$

Repeat forever:

$$V^{t+1} = TV^t$$

$$t \leftarrow t + 1$$

Q Iteration

Start from arbitrary an $Q^0 \in \mathbf{R}^{S \times A}$ and set $t = 0$
Based the purpose, set $T = T^\pi$ or $T = T^*$

Repeat forever:

$$Q^{t+1} = TQ^t$$

$$t \leftarrow t + 1$$